

Midterm & Concept Review

CS-585

Natural Language Processing

Sonjia Waxmonsky

Midterm Exam info

- Exam locations and times
 - Midterm: Thurs, Oct. 5, 8:35-9:50am
 - Room: SH-118 (Siegel 118)
- Midterm details
 - Timed: 70 minutes
 - Multiple-choice or similar (bubble answer sheet)
 - Closed book, closed note, no electronics
 - Bring pencils with erasers

Mid-Term Logistics

Remote exam-takers

- Arrangements should already be finalized
- Instructor cannot proctor via zoom or video

In-person exam-takes

- Exam will be timed to 70 minutes (for fairness with online sections); most students will not need entire time
- Please arrive by 8:35am to Siegel 118.
- Late arrivers will be promptly seated by teaching staff, exam cannot be extended
- Bring pencils with erasers! (Does not need to be #2)

Topics – Up to Midterm

	Content	Reading
1	Welcome, linguistic concepts	MS 3
2	Math review 1: probability and linear algebra	MS 2.1
3	Math review 2: information theory	MS 2.2
4	Words and pattern matching	MS 4.2, E-NLP 4.3
5	Lexical representations for NLP	MS 1.4
6	Neural nets 1: neural word embeddings	E-NLP 14
7	Word sense disambiguation	MS 7.1-7.3
8	Text categorization and naive Bayes	E-NLP 2.1-2.2, 4.4
9	Generalized linear models	E-NLP 2.5-2.6
10	Neural nets 2: feedforward networks	E-NLP 3.1-3.3.3
11	Sentiment analysis	E-NLP 4.1
12	Unsupervised methods in NLP	E-NLP 5.1

SAMPLE QUESTIONS

Sample Question - Linguistics

You are interested in studying how word meaning is influenced by suffix endings, such as –ed and –ing. Which fields are BEST suited to inform your investigation?

- A. Phonetics and phonology
- B. Morphology and semantics
- C. Historical Linguistics
- D. Image processing

Sample Question – Information Theory

What is measured with **cross-entropy**?

- A. The percentage of true positives that are detected in a binary classification model?
- B. The distance between two probability distributions
- C. The randomness or uncertainty of a random variable
- D. Any of the above

Sample Question – Regular expressions

Which string is matched by the following regular expression?

`\bhedge-?hog.\b`

- A. `hedge-hog`
- B. `hedge-hogs`
- C. `hedgehog`
- D. All of the above



classroomclipart.com
www.classroomclipart.com

Sample Question – Neural Networks

You are building a text categorization model, and your teammate shares they were able to **reduce model complexity** by **adding a layer** to their feed forward neural network. What is unusual about this statement?

- A. Neural networks cannot be feed-forward
- B. Adding a layer would generally increase complexity, not decrease complexity
- C. Complexity should always be increased, not decreased.
- D. All of the above

Sample Question – Metrics

You have built a rules-based text classifier that uses regular expressions to flag text messages that include a telephone number. This tool outputs 1 when it finds any string of 10 digits, and 0 otherwise. What is a good metric to evaluate this text classification tool? Select the BEST answer

- A. Precision
- B. ROC-AUC
- C. Cross-entropy
- D. All of the above

Review question

- Which of the following are NLP tasks?
- Which correspond to algorithms (or families of algorithms)? Which of those are unsupervised?
 - Word sense disambiguation
 - Sentiment analysis
 - Latent semantic analysis
 - word2vec