
Before we start:

HOMEWORK 1

Goals of Homework 1

- To apply the Information Theory concepts
- To begin to work with text data in python, and apply an open-source NLP package
 - NLTK
- To gain exposure to an openly available NLP research dataset
 - **GLUE Benchmark**

Homework 1 Steps

- Due Weds. September 13 at 11:59pm
- 100 points total
- 7 questions:
 - Download and prepare data
 - Write 4 functions and apply them to data
 - Answer 2 questions on Blackboard

GLUE Benchmark Datasets

- GLUE: General Language Understanding Evaluation
- Goal: "favor and encourage models that share general linguistic knowledge across tasks."
- 9 NLU tasks with a range of sources:
 - Classification: sentiment, linguistic acceptability,
 - Semantic similarity and equivalence
 - Inference: Question answering, entailment, pronoun resolution

<https://openreview.net/pdf?id=rJ4km2R5t7>

NLTK

- NLTK: Natural Language Toolkit
- Open-source python library for a range of NLP tasks, including Tokenization and Stemming

```
>>> word_tokenize("I wouldn't expect it to split this way!")
```

```
['I', 'would', "n't", 'expect', 'it', 'to', 'split', 'this', 'way', '!']
```



Homework 1 Help

- TA Office Hours start Tues Sept 5
 - Virtual and in-person
 - At least one office hour every weekday
 - Schedule posted in Blackboard – see "Content"
- Please bring python and machine setup questions to TA Office Hours
- Blackboard discussion group is available for questions and clarifications about instructions. See "HW 1 Discussion"

Homework 1 Setup

- **Start your Python setup early** if you are new to Python, or not yet set up on machine you will use for this class
- We recommend completing your setup **next week** (week of Sept 5-8). This means you can:
 1. Use python interpreter, e.g. `print("Hello World")`
 2. Unzip data files
 3. Run notebook "HW1_GettingStarted.ipynb"

Homework Submission

- Due Weds. September 13 at 11:59pm on Blackboard
- The homework late policy as posted in course syllabus applies to HW1
- Multiple submissions are allowed; no penalty for resubmission. Only last submission will be graded
- Submit early in case of last-minute technical issues

Words and Pattern Matching

CS-585

Natural Language Processing

Sonjia Waxmonsky

Regular expressions

- A formal language for specifying text strings
- How can we search for any of these?
 - woodchuck
 - woodchucks
 - Woodchuck
 - Woodchucks



Regular Expressions

In practice: Applied in Python `re` library

```
[1]: import re
```

```
[2]: my_animal = "TYPE_Wood_Chuck"  
     re.search("wood.*chuck",my_animal, re.IGNORECASE)
```

```
[2]: <re.Match object; span=(5, 15), match='Wood_Chuck'>
```

Regular Expressions: Disjunctions

- Letters inside square brackets []

Pattern	Matches
<code>[wW]oodchuck</code>	Woodchuck, woodchuck
<code>[1234567890]</code>	Any digit

- Ranges `[A-Z]`

Pattern	Matches	
<code>[A-Z]</code>	An upper case letter	<u>D</u> renched Blossoms
<code>[a-z]</code>	A lower case letter	<u>m</u> y beans were impatient
<code>[0-9]</code>	A single digit	Chapter <u>1</u> : Down the Rabbit Hole

Negation in Disjunction

- Negations `[^Ss]`
 - Caret means negation only when first in []

Pattern	Matches	
<code>[^A-Z]</code>	Not an upper case letter	O <u>y</u> fn pripetchik
<code>[^Ss]</code>	Neither 'S' nor 's'	<u>I</u> have no exquisite reason"
<code>[^e^]</code>	Neither e nor ^	Look he <u>r</u> e
<code>a^b</code>	The pattern a carat b	Look up <u>a^b</u> now

? * + .

Pattern		Matches
colou?r	Optional previous char	<u>color</u> <u>colour</u>
o*h!	0 or more of previous char	<u>h!</u> <u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
o+h!	1 or more of previous char	<u>oh!</u> <u>ooh!</u> <u>oooh!</u> <u>ooooh!</u>
baa+		<u>baa</u> <u>baaa</u> <u>baaaa</u> <u>baaaaa</u>
beg.n	Any char	<u>begin</u> <u>begun</u> <u>begun</u> <u>beg3n</u>

Anchors [^] \$

Pattern	Matches	
[^] [A-Z]	<u>P</u> alo Alto	Start of string
[^] [^A-Za-z]	<u>1</u> <u>"Hello"</u>	
\ . ^{\$}	The end <u>.</u>	End of string
. ^{\$}	The end <u>?</u> The end <u>!</u>	

Character classes

Pattern	Matches
<code>\s</code>	A whitespace character
<code>\S</code>	A non-whitespace character
<code>\d</code>	A digit (<code>[0-9]</code>)
<code>\D</code>	A non-digit
<code>\w</code>	A “word” character (<code>[0-9a-zA-Z_]</code>)
<code>\W</code>	A non-word character
<code>[:upper:]</code>	An upper-case letter
<code>[:lower:]</code>	A lower-case letter

Backreferences `(...) ...\n`

- Sometimes we want to know which part of the text matched a part of a pattern
- We can even use it within the pattern itself, by “capturing” it in parentheses

Pattern	Matches	
<code>(\d) [a-z] \1</code>	zsdfg <u>1a1</u> z213	A letter bracketed by the same number on each side
<code>^ (\d) (\d) .* \2\1\$</code>	<u>13awdfgasdf31</u>	A line starting with two digits, and ending with those two digits in reverse order

Example

- Find me all instances of the word “the” in a text.

the

Misses capitalized examples

[tT]he

Incorrectly returns other or theology

\b[tT]he\b

Summary

- Regular expressions are surprisingly important
 - Often the first model for any text processing
- For many tasks, we use machine learning
 - But regular expressions are used as features in the classifiers
 - Can be very useful in capturing generalizations

TOKENIZATION

Text Normalization

- Every NLP task needs to do text normalization:
 1. Segmenting/tokenizing words in running text
 2. Normalizing word formats
 3. Segmenting sentences in running text

How many words?

- I do uh main- mainly business data processing
 - Fragments, corrections, filled pauses
- Seuss's **cat** in the hat is different from other **cats**!
 - **Lemma**: same stem, part of speech, rough word sense
 - **cat** and **cats** = same lemma
 - **Wordform**: the full inflected surface form
 - **cat** and **cats** = different wordforms

How many words?

they lay back on the San Francisco grass and looked at the stars and their

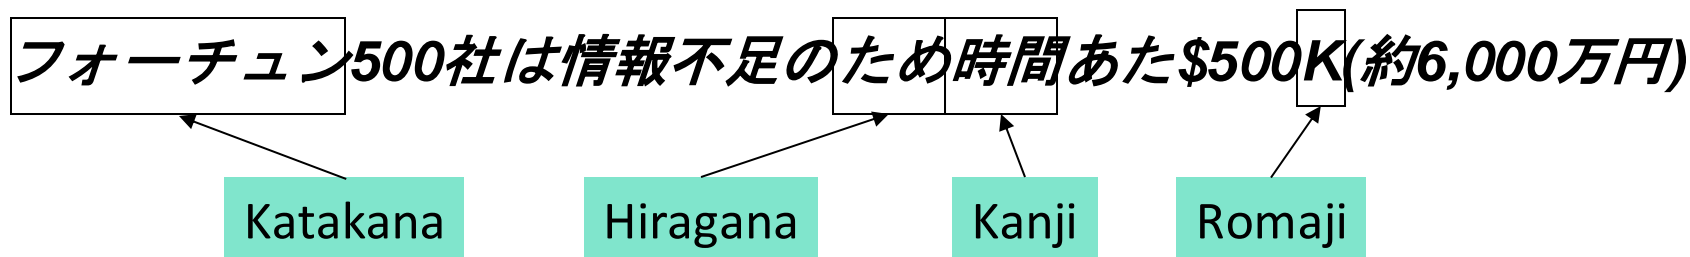
- **Type**: an element of the vocabulary.
- **Token**: an instance of that type in running text.
- How many? 15 tokens and 13 types
- Shakespeare - 31k types, 884k tokens

Tokenization: language issues

- French
 - *L'ensemble* → one token or two?
 - *L* ? *L'* ? *Le* ?
 - Want *l'ensemble* to match with *un ensemble*
- German noun compounds not segmented
 - *Lebensversicherungsgesellschaftsangestellter*
 - 'life insurance company employee'
 - German information retrieval needs **compound splitter**

Tokenization: language issues

- Chinese and Japanese -- no spaces between words:
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
 - Sharapova now lives in US southeastern Florida
- Japanese example: multiple alphabets; dates/amounts in multiple formats



Word Tokenization in Chinese

- Also called **Word Segmentation**
- Chinese words are composed of characters
 - Characters are generally 1 syllable and 1 morpheme.
 - Average word is 2.4 characters long.
- Standard baseline segmentation algorithm:
 - Maximum Matching (also called Greedy)

Maximum Matching Word Segmentation Algorithm (“greedy”)

Given a wordlist of Chinese, and a string.

- 1) Start a pointer at the beginning of the string
- 2) Find the longest word in dictionary that matches the string starting at pointer
- 3) Move the pointer over the word in string
- 4) Go to 2

Max-match segmentation illustration


- Thecatinthehat the cat in the hat
- Thetabledownthere the table down there
 theta bled own there

Doesn't generally work in English!

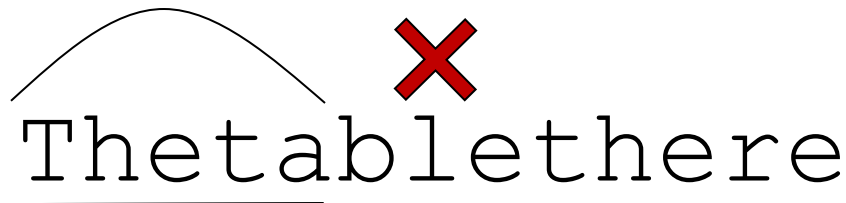
- But works astonishingly well in Chinese
 - 莎拉波娃现在居住在美国东南部的佛罗里达。
 - 莎拉波娃 现在 居住 在 美国 东南部 的 佛罗里达
- Modern probabilistic segmentation algorithms even better

Greedy matching


Thetabledownthere



Thetabl~~e~~there



Backtracking



Dynamic programming

Keep track of *intermediate results* (segments of string that can be parsed as a sequence of words)

Thetablethere

Location (character indices)	Parse
0-3	the
0-5	theta
0-8	the + table
0-11	the + table + the
0-13	the + table + there

WORD NORMALIZATION & STEMMING

Normalization

- Need to “normalize” terms
 - Information Retrieval: indexed text & query terms must have same form.
 - We want to match ***U.S.A.*** and ***USA***
- We implicitly define equivalence classes of terms
 - e.g., deleting periods in a term
- Alternative: asymmetric expansion:
 - Enter: ***window*** Search: ***window, windows***
 - Enter: ***windows*** Search: ***Windows, windows, window***
 - Enter: ***Windows*** Search: ***Windows***
- Potentially more powerful, but less efficient

Case folding

- Applications like Information Retrieval: reduce all letters to lower case
 - Since users tend to use lower case
 - Possible exception: upper case in mid-sentence?
 - e.g., **General Motors**
 - **Fed** vs. **fed**
 - **SAT** vs. **sat**
- For sentiment analysis, MT, Information extraction
 - Case is helpful (**US** versus **us** is important)

Lemmatization

- Reduce inflections or variant forms to **base** form
 - *am, are, is* → *be*
 - *car, cars, car's, cars'* → *car*
- *the boy's cars are different colors* → *the boy car be different color*
- Lemmatization: have to find correct dictionary headword form
- Machine translation
 - Spanish **quiero** ('I want'), **quieres** ('you want') same lemma as **querer** 'want'

Morphology

- **Morphemes:**

- The small meaningful units that make up words
- **Stems**: The core meaning-bearing units
- **Affixes**: Bits and pieces that adhere to stems
 - Often with grammatical functions

Stemming

- Reduce terms to their stems
- *Stemming* is crude chopping of affixes
 - language dependent
 - e.g., ***automate(s), automatic, automation*** all reduced to ***automat.***

*for example compressed
and compression are both
accepted as equivalent to
compress.*



for exampl compress and
compress ar both accept
as equival to compress

Complex morphology

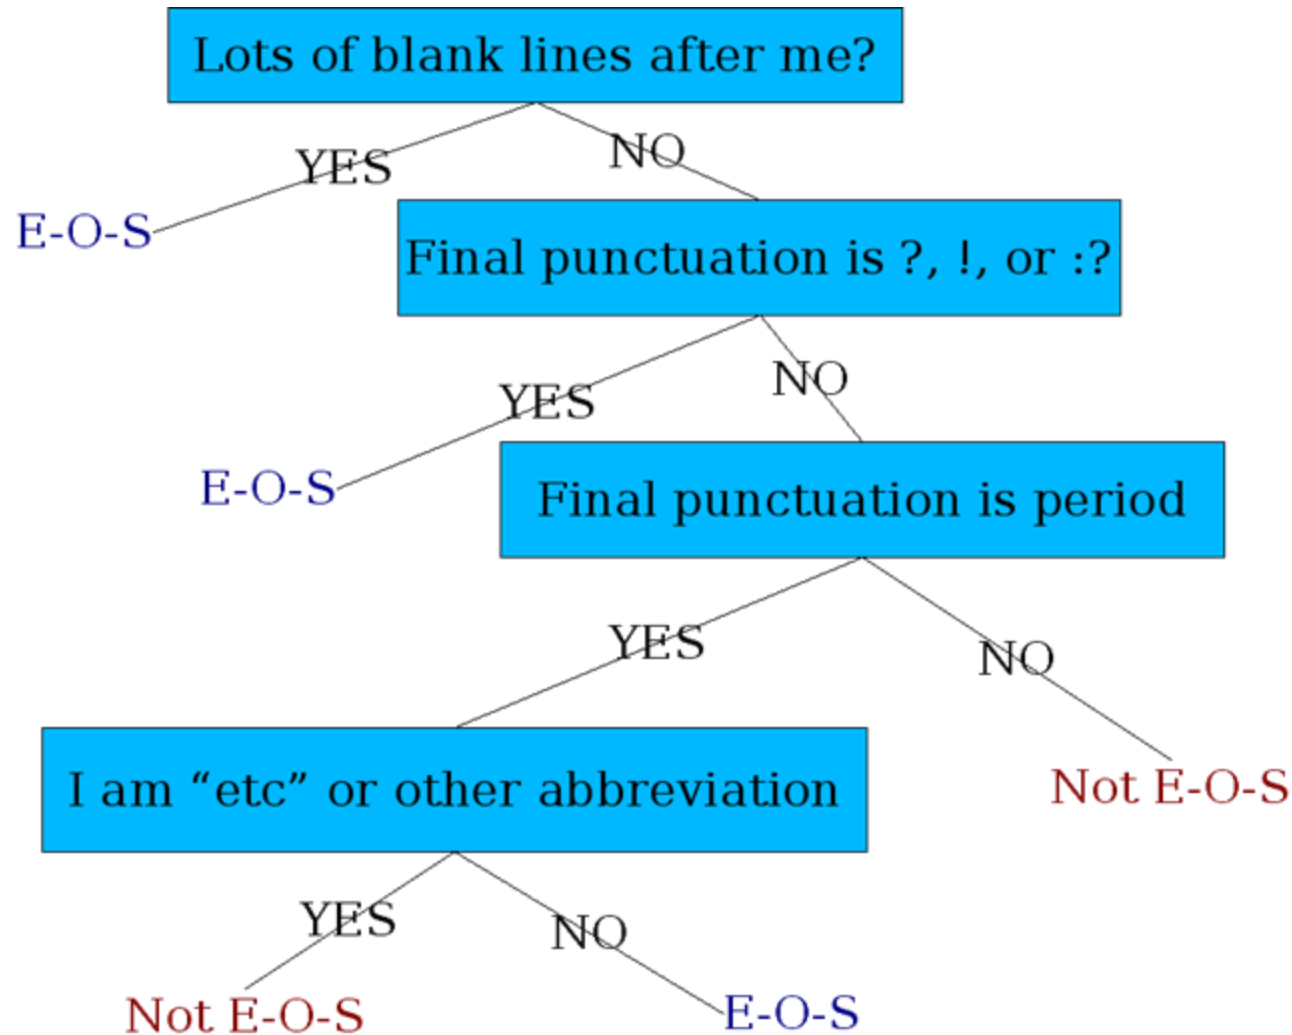
- Some languages require complex morpheme segmentation
 - Turkish
 - **Uygarlaştıramadıklarımızdanmışsınızcasına**
 - '(behaving) as if you are among those whom we could not civilize'
 - **Uygar** 'civilized' + **laş** 'become'
 - + **tır** 'cause' + **ama** 'not able'
 - + **dik** 'past' + **lar** 'plural'
 - + **ımız** 'p1pl' + **dan** 'abl'
 - + **mış** 'past' + **sınız** '2pl' + **casına** 'as if'

SENTENCE SEGMENTATION

Where to break sentences?

- !, ? are relatively unambiguous
- Period “.” is very ambiguous
 - Sentence boundary
 - Abbreviations like Inc. or Dr.
 - Numbers like .02% or 4.3
- Build a classifier
 - Looks at a “.”
 - Decides EndOfSentence/NotEndOfSentence
 - Classifiers: hand-written rules, regular expressions, or machine-learning

A Decision Tree



More sophisticated features

- Case of word preceding “.”:
Upper, Lower, Cap, Number
- Case of word following “.”:
Upper, Lower, Cap, Number
- Numeric features
 - Length of word preceding “.”
 - Probability (word preceding “.” occurs at end-of-sent)
 - Probability (word after “.” occurs at beginning-of-s)

Implementing Decision Trees

- A decision tree is just an if-then-else statement
- The interesting question is choosing the features
- Setting up the structure is often too hard to do by hand
 - Only possible for very simple features, domains
 - For numeric features, it's too hard to pick each threshold
 - Instead, structure usually learned by machine learning from a training corpus (later in the course)

HOMEWORK 1

Goals of Homework 1

- To apply the Information Theory concepts
- To begin to work with text data in python, and apply an open-source NLP package
 - NLTK
- To gain exposure to an openly available NLP research dataset
 - **GLUE Benchmark**

Homework 1 Steps

- Due Weds. September 13 at 11:59pm
- 100 points total
- 7 questions:
 - Download and prepare data
 - Write 4 functions and apply them to data
 - Answer 2 questions on Blackboard

GLUE Benchmark Datasets

- GLUE: General Language Understanding Evaluation
- Goal: "favor and encourage models that share general linguistic knowledge across tasks."
- 9 NLU tasks with a range of sources:
 - Classification: sentiment, linguistic acceptability,
 - Semantic similarity and equivalence
 - Inference: Question answering, entailment, pronoun resolution

<https://openreview.net/pdf?id=rJ4km2R5t7>

NLTK

- NLTK: Natural Language Toolkit
- Open-source python library for a range of NLP tasks, including Tokenization and Stemming

```
>>> word_tokenize("I wouldn't expect it to split this way!")
```

```
['I', 'would', "n't", 'expect', 'it', 'to', 'split', 'this', 'way', '!']
```



Homework 1 Help

- TA Office Hours start Tues Sept 5
 - Virtual and in-person
 - At least one office hour every weekday
 - Schedule posted in Blackboard – see "Content"
- Please bring python and machine setup questions to TA Office Hours
- Blackboard discussion group is available for questions and clarifications about instructions. See "HW 1 Discussion"

Homework 1 Setup

- **Start your Python setup early** if you are new to Python, or not yet set up on machine you will use for this class
- We recommend completing your setup **next week** (week of Sept 5-8). This means you can:
 1. Use python interpreter, e.g. `print("Hello World")`
 2. Unzip data files
 3. Run notebook "HW1_GettingStarted.ipynb"

Homework Submission

- Due Weds. September 13 at 11:59pm on Blackboard
- The homework late policy as posted in course syllabus applies to HW1
- Multiple submissions are allowed; no penalty for resubmission. Only last submission will be graded
- Submit early in case of last-minute technical issues